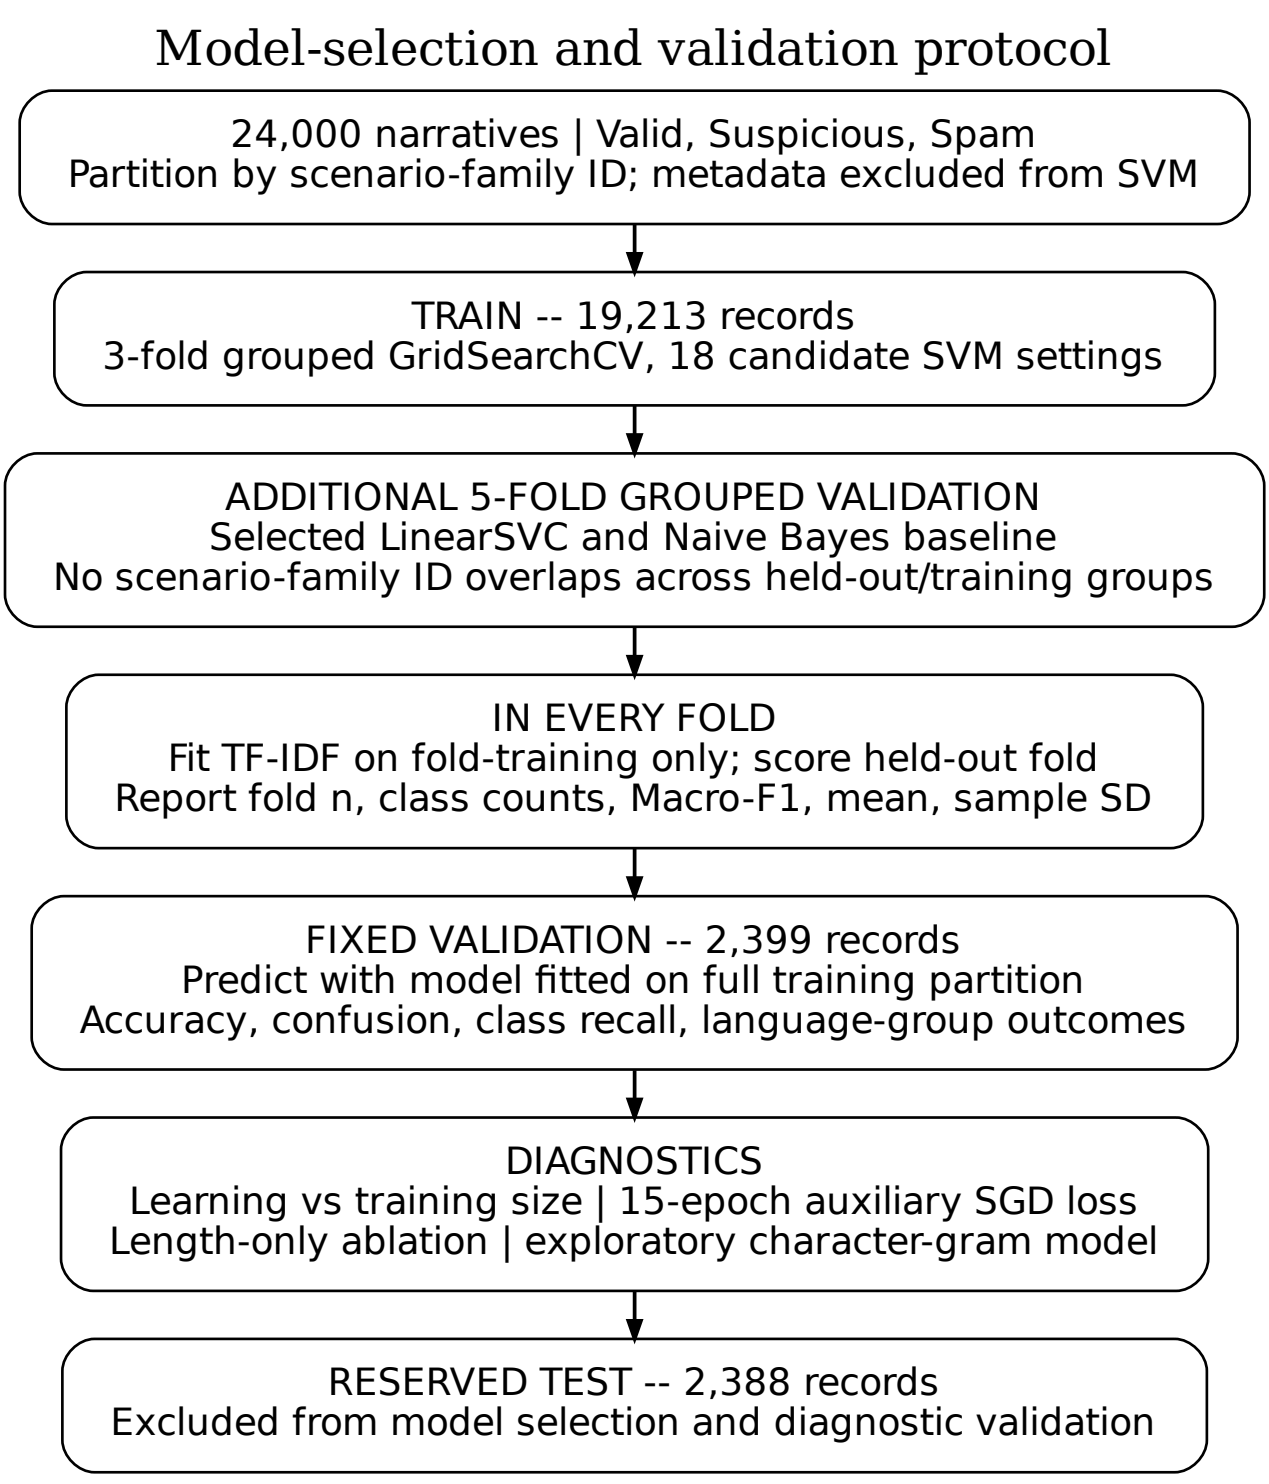
Model-selection and validation protocol
24,000 narratives | Valid, Suspicious, Spam
Partition by scenario-family ID; metadata excluded from SVM
TRAIN -- 19,213 records
3-fold grouped GridSearchCV, 18 candidate SVM settings
ADDITIONAL 5-FOLD GROUPED VALIDATION
Selected LinearSVC and Naive Bayes baseline
No scenario-family ID overlaps across held-out/training groups
IN EVERY FOLD
Fit TF-IDF on fold-training only; score held-out fold
Report fold n, class counts, Macro-F1, mean, sample SD
FIXED VALIDATION -- 2,399 records
Predict with model fitted on full training partition
Accuracy, confusion, class recall, language-group outcomes
DIAGNOSTICS
Learning vs training size | 15-epoch auxiliary SGD loss
Length-only ablation | exploratory character-gram model
RESERVED TEST -- 2,388 records
Excluded from model selection and diagnostic validation