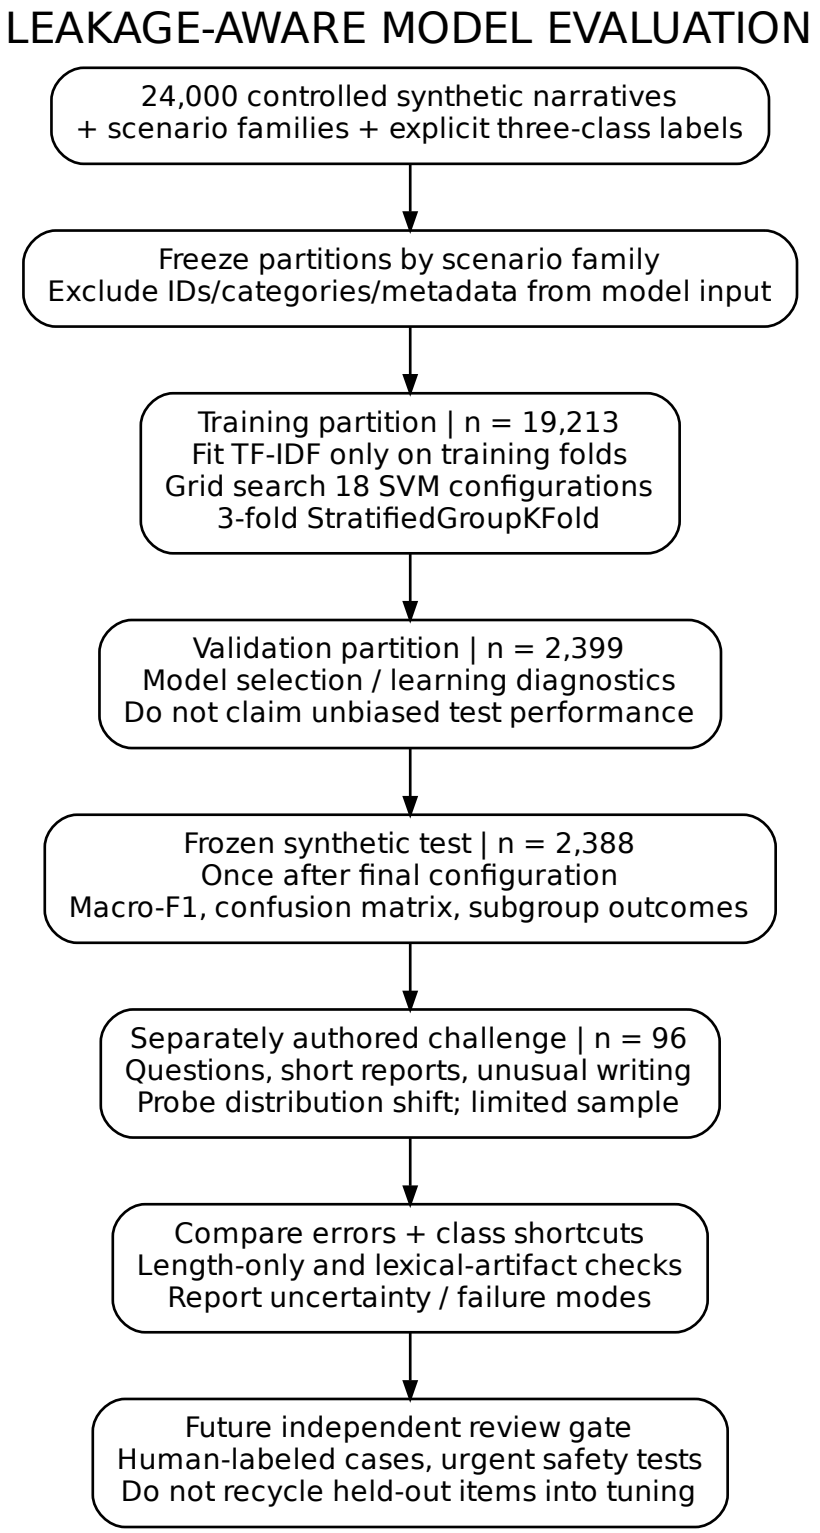
LEAKAGE-AWARE MODEL EVALUATION
24,000 controlled synthetic narratives
+ scenario families + explicit three-class labels
Freeze partitions by scenario family
Exclude IDs/categories/metadata from model input
Training partition | n = 19,213
Fit TF-IDF only on training folds
Grid search 18 SVM configurations
3-fold StratifiedGroupKFold
Validation partition | n = 2,399
Model selection / learning diagnostics
Do not claim unbiased test performance
Frozen synthetic test | n = 2,388
Once after final configuration
Macro-F1, confusion matrix, subgroup outcomes
Separately authored challenge | n = 96
Questions, short reports, unusual writing
Probe distribution shift; limited sample
Compare errors + class shortcuts
Length-only and lexical-artifact checks
Report uncertainty / failure modes
Future independent review gate
Human-labeled cases, urgent safety tests
Do not recycle held-out items into tuning