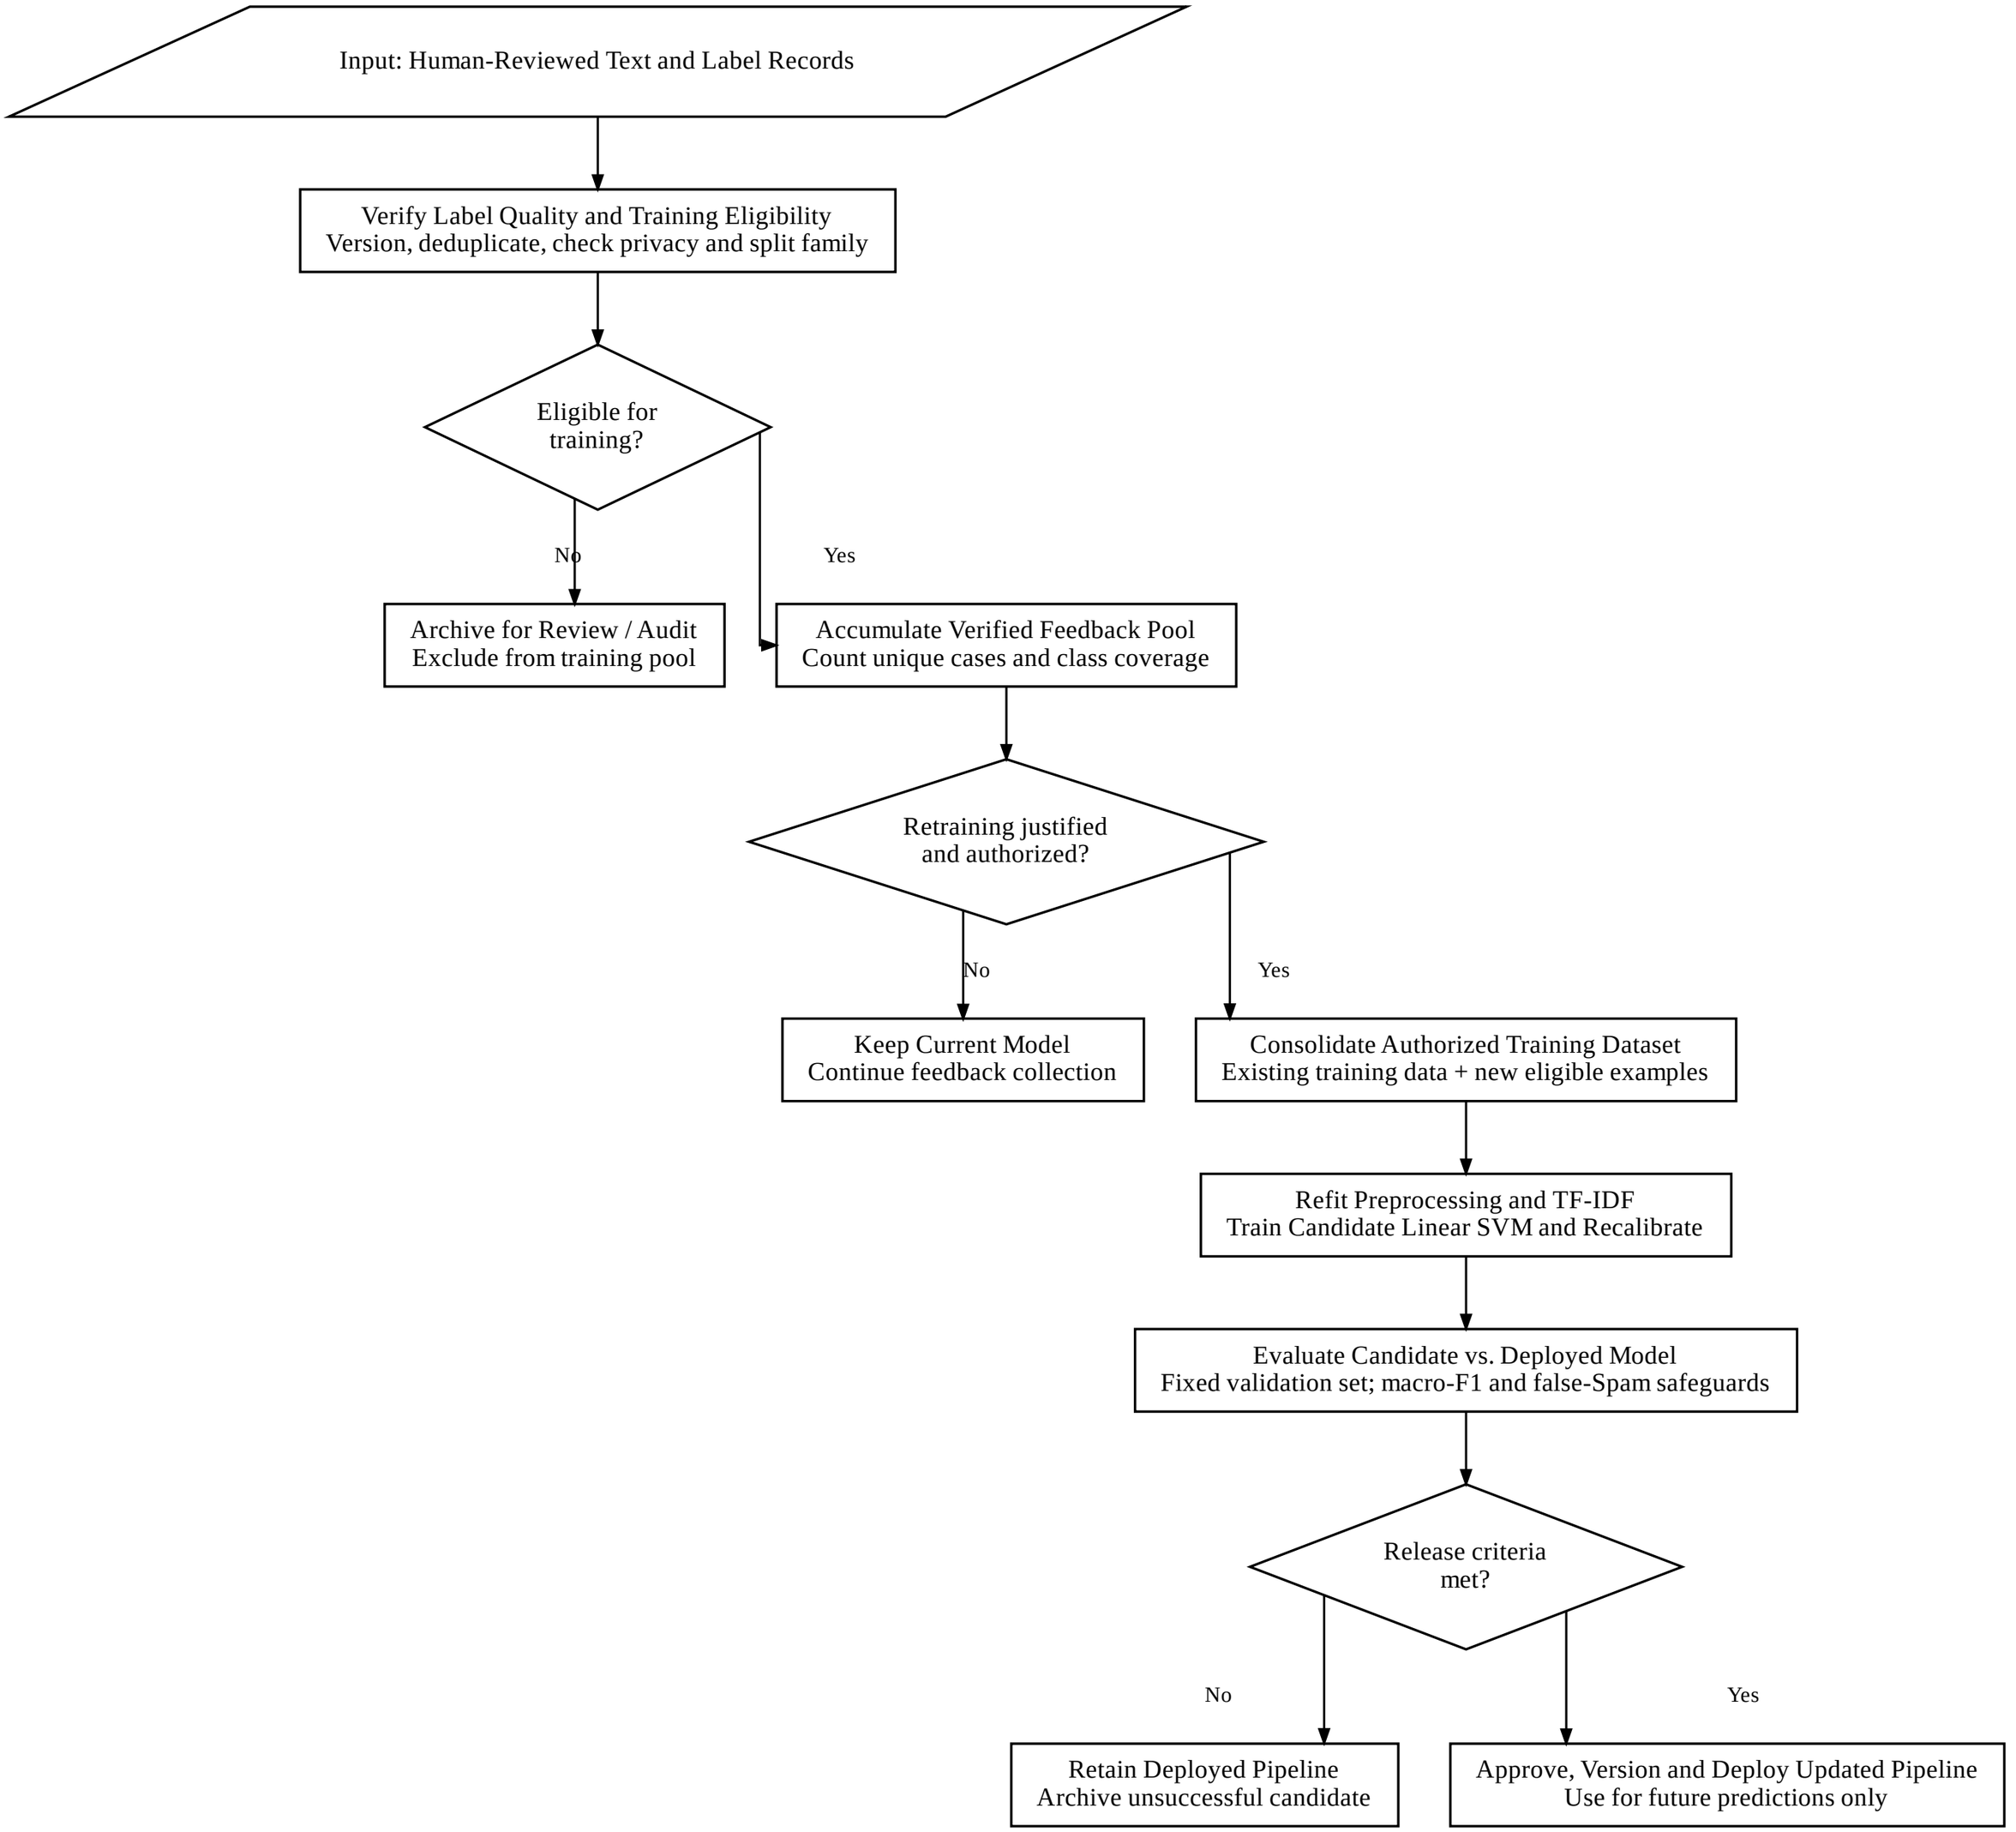
Input: Human-Reviewed Text and Label Records
Verify Label Quality and Training Eligibility
Version, deduplicate, check privacy and split family
Eligible for training?
No
Yes
Archive for Review / Audit
Exclude from training pool
Accumulate Verified Feedback Pool
Count unique cases and class coverage
Retraining justified and authorized?
No
Yes
Keep Current Model
Continue feedback collection
Consolidate Authorized Training Dataset
Existing training data + new eligible examples
Refit Preprocessing and TF-IDF
Train Candidate Linear SVM and Recalibrate
Evaluate Candidate vs. Deployed Model
Fixed validation set; macro-F1 and false-Spam safeguards
Release criteria met?
No
Yes
Retain Deployed Pipeline
Archive unsuccessful candidate
Approve, Version and Deploy Updated Pipeline
Use for future predictions only